

Data wrangling with **dplyr**

STAT 220

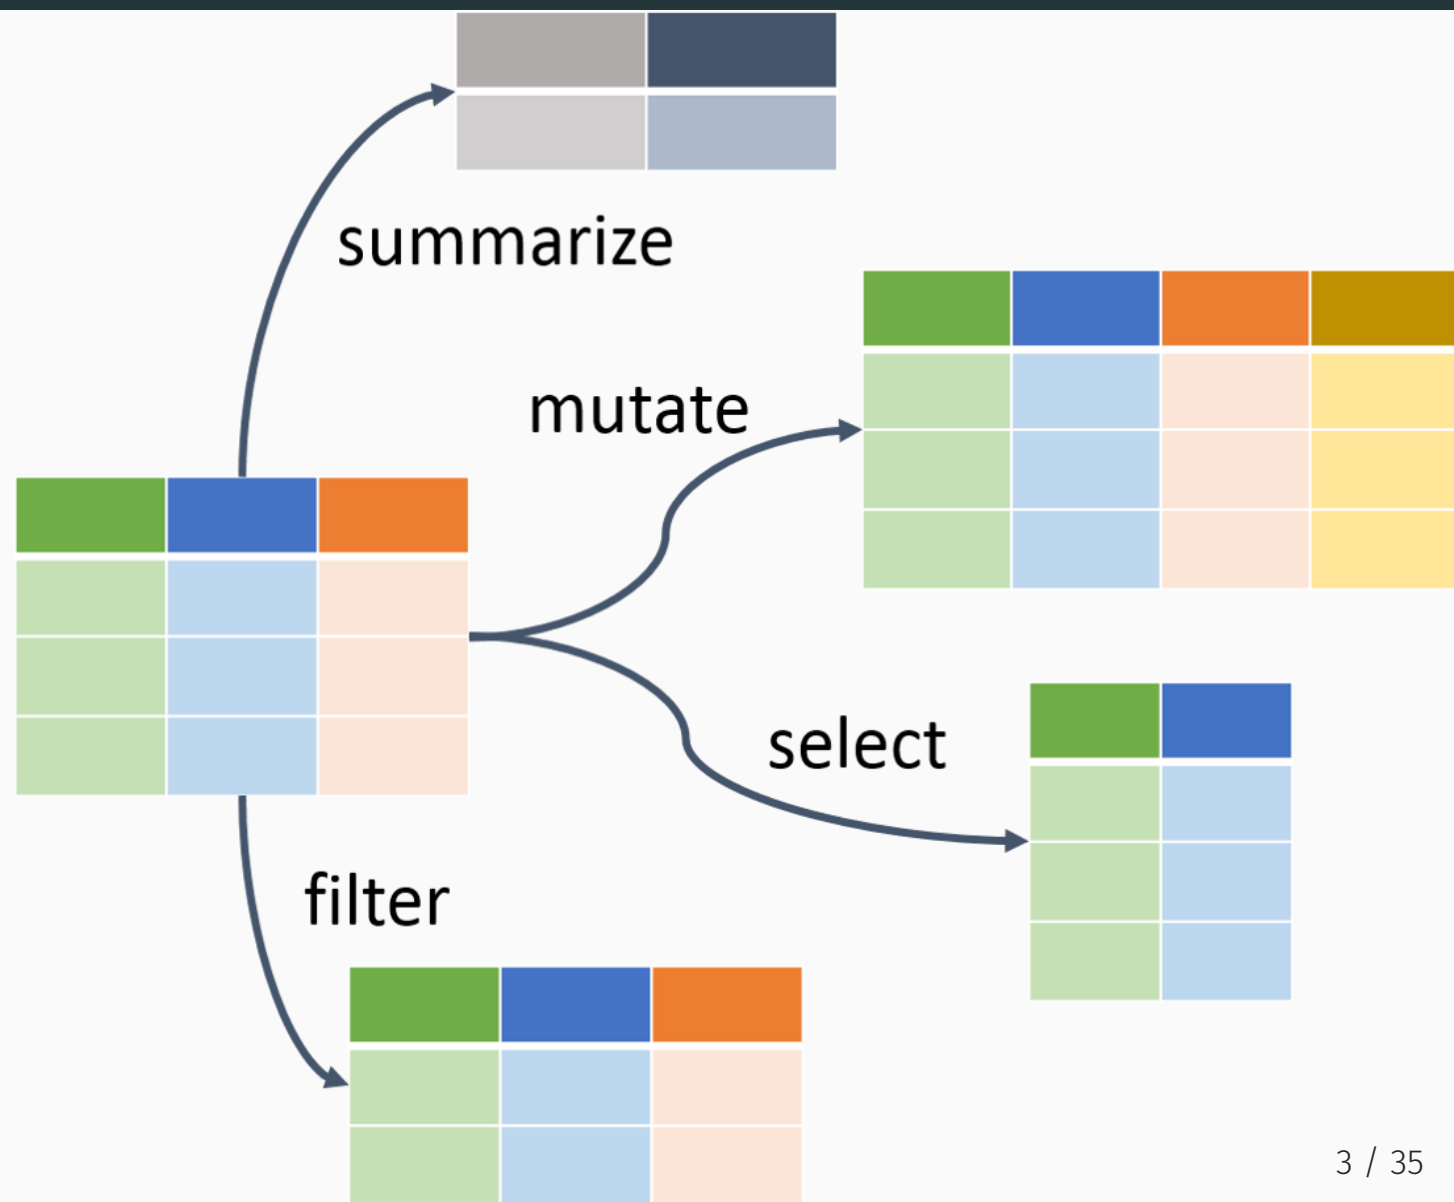

Bastola

January 17 2022

Data Wrangling

the process of cleaning and unifying messy and complex data sets for easy access and analysis

- "data janitor work"
- importing data
- cleaning data
- changing shape of data
- fixing errors and poorly formatted data elements
- transforming columns and rows
- filtering, subsetting



Essential data wrangling tasks

1. Extract a subset of columns
2. Extract a subset of rows
3. Order the rows from smallest to largest (or largest to smallest)
4. Compute a table of summary statistics, perhaps by group
5. Create new columns

The Five Verbs

The creator of `dplyr`, [Hadley Wickham](#), argues that most of the operations on a data table can be achieved with

- *select()*
- *filter()*
- *mutate()*
- *arrange()*
- *summarize()*

Babynames Dataset

- Names of babies born in the U.S. between 1880 and 2017
- Source: Social Security Administration

```
library(babynames)
glimpse(babynames)
Rows: 1,924,665
Columns: 5
$ year <dbl> 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880,...
$ sex  <chr> "F", "F", "F", "F", "F", "F", "F", "F", "F", "F", "F", "F", "F", ...
$ name <chr> "Mary", "Anna", "Emma", "Elizabeth", "Minnie", "Margaret", "Ida", ...
$ n    <int> 7065, 2604, 2003, 1939, 1746, 1578, 1472, 1414, 1320, 1288, 1258, ...
$ prop <dbl> 0.07238359, 0.02667896, 0.02052149, 0.01986579, 0.01788843, 0.016...
```

Find a subset of the columns using `select()`:

- `select()`: take a subset of the columns (variables/features)

```
babynames %>%  
  select(year, name, n) %>%  
  head()
```

```
# A tibble: 6 × 3
```

	year	name	n
	<dbl>	<chr>	<int>
1	1880	Mary	7065
2	1880	Anna	2604
3	1880	Emma	2003
4	1880	Elizabeth	1939
5	1880	Minnie	1746
6	1880	Margaret	1578

Using %>%

- `%>%` passes result on left into first argument of function on right
- Chaining functions together lets you read Left-to-right, top-to-bottom

Using %>%

You can build up a series of pipes

```
babynames %>%                                # dataframe first and then...  
  select(year, name, n) %>%  
  head()
```

Using %>%

You can build up a series of pipes

```
babynames %>%                                # dataframe first and then...  
  select(year, name, n) %>%                  # select columns year, name, and n  
  head()
```

Using %>%

You can build up a series of pipes

```
babynames %>%           # dataframe first and then...  
  select(year, name, n) %>% # select columns year, name, and n  
  head()                # display header of the data frame
```

A tibble: 6 × 3

	year	name	n
	<dbl>	<chr>	<int>
1	1880	Mary	7065
2	1880	Anna	2604
3	1880	Emma	2003
4	1880	Elizabeth	1939
5	1880	Minnie	1746
6	1880	Margaret	1578

select() helpers

: select range of columns

```
select(gapminder, income:population)
```

- select every column but

```
select(gapminder, -c(income,population))
```

select() helpers

`starts_with()` select columns that start with...

```
select(gapminder, starts_with("p"))
```

`ends_with()` select columns that end with...

```
select(gapminder, ends_with("y"))
```

`contains()` select columns whose names contain...

```
select(gapminder, contains("e"))
```

Your Turn 1

- Please git clone the repository on [data wrangling activity](#) from the course GitHub organization.
- Which of these is **NOT** a way to select the `name` and `n` columns together?

```
select(babynames, -c(year, sex, prop))
```

```
select(babynames, name:n)
```

```
select(babynames, starts_with("n"))
```

```
select(babynames, ends_with("n"))
```

02:00

Find a subset of the rows using *filter()*:

- `filter()`: take a subset of the rows (observations)

```
babynames %>%  
  filter(name == "Bella") %>%  
  head()  
# A tibble: 6 × 5  
   year sex  name    n  prop  
  <dbl> <chr> <chr> <int> <dbl>  
1  1880 F    Bella    13 0.000133  
2  1881 F    Bella    24 0.000243  
3  1882 F    Bella    16 0.000138  
4  1883 F    Bella    17 0.000142  
5  1884 F    Bella    31 0.000225  
6  1885 F    Bella    25 0.000176
```

Use both *filter()* and *select()*

```
bella <- babynames %>%  
  filter(name == "Bella") %>%  
  select(year, name, sex, n)
```

```
head(bella)  
# A tibble: 6 × 4  
   year name  sex      n  
  <dbl> <chr> <chr> <int>  
1  1880 Bella  F      13  
2  1881 Bella  F      24  
3  1882 Bella  F      16  
4  1883 Bella  F      17  
5  1884 Bella  F      31  
6  1885 Bella  F      25
```

```
dim(bella)  
[1] 144  4  
class(bella)  
[1] "tbl_df"      "tbl"        "data.frame"
```


Some Operators

Operator	Definition
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to
!=	not equal to
x & y	x AND y
x %in% y	test if x is in y

Your Turn 2

Use `filter()` with the logical operators to extract:

- All of the names where prop is greater than or equal to 0.08
- All of the babies named "Rose"
- All of the names that have a missing value for n

02:00

summarize() or summarise()

If we want to compare summary statistics, we might use `summarize()`.

```
babynames %>%  
  filter(name == "Bella", sex == "F") %>%  
  summarise(total = sum(n), max = max(n), mean = mean(n))  
# A tibble: 1 × 3  
  total    max  mean  
  <int> <int> <dbl>  
1  57411  5121  416.
```

```
babynames %>%  
  filter(name == "Bella", sex == "F") %>%  
  summarize(n = n())  
# A tibble: 1 × 1  
      n  
  <int>  
1   138
```

```
babynames %>%  
  summarize(nname = n_distinct(name))  
# A tibble: 1 × 1  
  nname  
  <int>  
1  97310
```

Your Turn 3


Use the codes mentioned so far to compute three statistics:

1. the total number of children who ever had your name
2. the maximum number of children given your name in a single year
3. the mean number of children given your name per year

04:00

Boolean operators

For help, `?base::Logic`



Source: *R for Data Science*, by Golemund & Wickham

Your Turn 4

Use Boolean operators to alter the code below to return only the rows that contain:

- Girls named Rhea
- Names that were used by exactly 5 or 6 children in 1990
- Names that are one of Apple, Yoroi, Ada

```
filter(babynames, name == "Ada" | name == "Rhea")
```

04:00

Using `group_by()`

```
babynames %>%
```

```
  group_by(year, sex)
```

```
# A tibble: 1,924,665 × 5
```

```
# Groups:   year, sex [276]
```

	year	sex	name	n	prop
	<dbl>	<chr>	<chr>	<int>	<dbl>
1	1880	F	Mary	7065	0.0724
2	1880	F	Anna	2604	0.0267
3	1880	F	Emma	2003	0.0205
4	1880	F	Elizabeth	1939	0.0199
5	1880	F	Minnie	1746	0.0179
6	1880	F	Margaret	1578	0.0162
7	1880	F	Ida	1472	0.0151
8	1880	F	Alice	1414	0.0145
9	1880	F	Bertha	1320	0.0135
10	1880	F	Sarah	1288	0.0132

```
# ... with 1,924,655 more rows
```

Using `group_by()` along with `summarize()`

```
babynames %>%  
  group_by(year) %>%  
  summarise(total = sum(n))
```

```
# A tibble: 138 × 2
```

	year	total
	<dbl>	<int>
1	1880	201484
2	1881	192696
3	1882	221533
4	1883	216946
5	1884	243462
6	1885	240854
7	1886	255317
8	1887	247394
9	1888	299473
10	1889	288946

```
# ... with 128 more rows
```


mutate()

- `mutate()` lets us create new variables based on manipulations of the old variables

```
babynames <- babynames %>%  
  group_by(year) %>%  
  mutate(percent = prop * 100)  
head(babynames)  
# A tibble: 6 × 6  
# Groups:   year [1]  
   year sex  name      n  prop percent  
   <dbl> <chr> <chr>   <int> <dbl>   <dbl>  
1  1880 F    Mary    7065 0.0724    7.24  
2  1880 F    Anna    2604 0.0267    2.67  
3  1880 F    Emma    2003 0.0205    2.05  
4  1880 F  Elizabeth 1939 0.0199    1.99  
5  1880 F   Minnie   1746 0.0179    1.79  
6  1880 F  Margaret 1578 0.0162    1.62
```

arrange()

Order rows from smallest to largest

```
arrange(.data, ...)
```

```
arrange(babynames, n)
```

desc()

Changes ordering from largest to smallest.

```
arrange(babynames, desc(n))
```

```
# A tibble: 1,924,665 × 6
```

```
# Groups:   year [138]
```

	year	sex	name	n	prop	percent
	<dbl>	<chr>	<chr>	<int>	<dbl>	<dbl>
1	1947	F	Linda	99686	0.0548	5.48
2	1948	F	Linda	96209	0.0552	5.52
3	1947	M	James	94756	0.0510	5.10
4	1957	M	Michael	92695	0.0424	4.24
5	1947	M	Robert	91642	0.0493	4.93
6	1949	F	Linda	91016	0.0518	5.18
7	1956	M	Michael	90620	0.0423	4.23
8	1958	M	Michael	90520	0.0420	4.20
9	1948	M	James	88588	0.0497	4.97
10	1954	M	Michael	88514	0.0428	4.28

```
# ... with 1,924,655 more rows
```

Most common names in 1990

```
babynames %>%
```

```
  filter(year == 1990) %>%
```

```
  arrange(desc(prop))
```

```
# A tibble: 24,719 × 6
```

```
# Groups:   year [1]
```

	year	sex	name	n	prop	percent
	<dbl>	<chr>	<chr>	<int>	<dbl>	<dbl>
1	1990	M	Michael	65282	0.0303	3.03
2	1990	M	Christopher	52332	0.0243	2.43
3	1990	F	Jessica	46475	0.0226	2.26
4	1990	F	Ashley	45558	0.0222	2.22
5	1990	M	Matthew	44800	0.0208	2.08
6	1990	M	Joshua	43216	0.0201	2.01
7	1990	F	Brittany	36538	0.0178	1.78
8	1990	F	Amanda	34408	0.0168	1.68
9	1990	M	Daniel	33815	0.0157	1.57
10	1990	M	David	33742	0.0157	1.57

```
# ... with 24,709 more rows
```

top_n()

Most common name in each year

```
babynames %>%
```

```
  group_by(year) %>%
```

```
  top_n(1, prop)
```


```
# A tibble: 138 × 6
```

```
# Groups:   year [138]
```

	year	sex	name	n	prop	percent
	<dbl>	<chr>	<chr>	<int>	<dbl>	<dbl>
1	1880	M	John	9655	0.0815	8.15
2	1881	M	John	8769	0.0810	8.10
3	1882	M	John	9557	0.0783	7.83
4	1883	M	John	8894	0.0791	7.91
5	1884	M	John	9388	0.0765	7.65
6	1885	M	John	8756	0.0755	7.55
7	1886	M	John	9026	0.0758	7.58
8	1887	M	John	8110	0.0742	7.42
9	1888	M	John	9247	0.0712	7.12

Your Turn 5

Find the year your first name is most common.



03:00

Vectorized function

a function that takes a vector as input, is applied to every element of the vector, and returns a vector (of the original length) as output

min_rank()

A go to ranking function (ties share the lowest rank)

```
min_rank(c(50, 100, 1000))  
[1] 1 2 3
```

```
min_rank(desc(c(50, 100, 1000)))  
[1] 3 2 1
```


Your Turn 6

- Use `min_rank()` and `mutate()` to rank each row in `babynames` from largest `prop` to smallest `prop`.
- Compute each name's rank within its `year` and `sex`.
- Then compute the median rank for each combination of `name` and `sex`, and arrange the results from highest median rank to lowest.

06 : 00

Summary

- Extract variables with `select()`
- Extract cases with `filter()`
- Arrange cases, with `arrange()`
- Make tables of summaries with `summarize()`
- Make new variables, with `mutate()`

Acknowledgement: some of the slides are based on previous works of Adam Loy and Katie St. Clair.