

More Data Wrangling

Fall 2022

September 26 2022

rename()

```
pollution <- tribble(
  ~city, ~size, ~amount,
  "New York", "large", 23.4,
  "New York", "small", 14.5,
  "London", "large", 22.2,
  "London", "small", 16.7,
  "Beijing", "large", 121.0,
  "Beijing", "small", 56.9
)
```

```
pollution %>% rename(quantity = amount)
```

```
# A tibble: 6 × 3
```

	city	size	quantity
	<chr>	<chr>	<dbl>
1	New York	large	23.4
2	New York	small	14.5
3	London	large	22.2
4	London	small	16.7
5	Beijing	large	121
6	Beijing	small	56.9

Mutating multiple columns at once: `mutate_*`

- variants of `mutate()` that are useful for mutating multiple columns at once
 - `mutate_at()`, `mutate_if()`, `mutate_all()`, etc.
- which columns get mutated depends on a predicate, can be:
 - a function that returns TRUE/FALSE like `is.numeric()`, or
 - variable names through `vars()`

```
pollution %>%
  mutate_at(vars(city:amount), toupper)
# A tibble: 6 × 3
  city      size amount
  <chr>    <chr> <chr>
1 NEW YORK LARGE 23.4
2 NEW YORK SMALL 14.5
3 LONDON   LARGE 22.2
4 LONDON   SMALL 16.7
5 BEIJING  LARGE 121
6 BEIJING  SMALL 56.9
```

```
pollution %>%
  mutate_if(is.double, round, digits = 0)
# A tibble: 6 × 3
  city      size amount
  <chr>    <chr> <dbl>
1 New York large    23
2 New York small    14
3 London   large    22
4 London   small    17
5 Beijing  large   121
6 Beijing  small    57
```

Selecting & renaming multiple columns


- `select_*`() & `rename_*`() are variants of `select()` and `rename()`
- use like `mutate_*`() options on previous slide

What do these commands do?

```
pollution %>% select_if(is.numeric)
pollution %>% rename_all(toupper)
pollution %>% rename_if(is.character, toupper)
pollution %>% rename_at(vars(contains("it")), toupper)
```

Boolean operators

For help, `?base::Logic`



Group Activity 1


05:00



- Let's go over to maize server/ local Rstudio and our class [moodle](#)
- Get the class activity 7.Rmd file
- Work on problems 1-2
- Ask me questions

Two-table verbs

- `inner_join()` - Merge two datasets. Exclude all unmatched rows.
- `full_join()` - Merge two datasets. Keep all observations.
- `left_join()` - Merge two datasets. Keep all observations from the origin table.
- `right_join()` - Merge two datasets. Keep all observations from the destination table.
- `anti_join()` - Drops all observations in origin that have a match in destination table.



Mutating Joins

- `left_join()`
- `right_join()`
- `inner_join()`
- `full_join()`

Differ in their behavior when a match is not found

Flights data

```
library(nycflights13)
flights2 <- flights %>%
  select(year:day, hour, origin, dest, tailnum, carrier)
```

```
head(flights2)
# A tibble: 6 × 8
   year month   day   hour origin dest  tailnum carrier
  <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>
1  2013     1     1     5 EWR    IAH   N14228   UA
2  2013     1     1     5 LGA    IAH   N24211   UA
3  2013     1     1     5 JFK    MIA   N619AA   AA
4  2013     1     1     5 JFK    BQN   N804JB   B6
5  2013     1     1     6 LGA    ATL   N668DN   DL
6  2013     1     1     5 EWR    ORD   N39463   UA
```

Airline information

```
head(airlines)
# A tibble: 6 × 2
  carrier name
  <chr>    <chr>
1 9E      Endeavor Air Inc.
2 AA      American Airlines Inc.
3 AS      Alaska Airlines Inc.
4 B6      JetBlue Airways
5 DL      Delta Air Lines Inc.
6 EV      ExpressJet Airlines Inc.
```

left_join()

```
flights2 %>%  
  left_join(airlines)  
# A tibble: 336,776 × 9  
   year month   day hour origin dest tailnum carrier name  
   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr>  
1  2013     1     1     5 EWR   IAH  N14228  UA    United Air Lines Inc.  
2  2013     1     1     5 LGA   IAH  N24211  UA    United Air Lines Inc.  
3  2013     1     1     5 JFK   MIA  N619AA  AA    American Airlines Inc.  
4  2013     1     1     5 JFK   BQN  N804JB  B6    JetBlue Airways  
5  2013     1     1     6 LGA   ATL  N668DN  DL    Delta Air Lines Inc.  
6  2013     1     1     5 EWR   ORD  N39463  UA    United Air Lines Inc.  
7  2013     1     1     6 EWR   FLL  N516JB  B6    JetBlue Airways  
8  2013     1     1     6 LGA   IAD  N829AS  EV    ExpressJet Airlines Inc.  
9  2013     1     1     6 JFK   MCO  N593JB  B6    JetBlue Airways  
10 2013     1     1     6 LGA   ORD  N3ALAA  AA    American Airlines Inc.  
# ... with 336,766 more rows
```

Planes information

```
head(planes)
# A tibble: 6 × 9
  tailnum  year type      manif...1 model engines seats speed engine
  <chr>    <int> <chr>      <chr>    <chr>   <int> <int> <int> <chr>
1 N10156   2004 Fixed wing multi engine EMBRAER EMB-...     2    55    NA Turbo...
2 N102UW   1998 Fixed wing multi engine AIRBUS... A320...     2   182    NA Turbo...
3 N103US   1999 Fixed wing multi engine AIRBUS... A320...     2   182    NA Turbo...
4 N104UW   1999 Fixed wing multi engine AIRBUS... A320...     2   182    NA Turbo...
5 N10575   2002 Fixed wing multi engine EMBRAER EMB-...     2    55    NA Turbo...
6 N105UW   1999 Fixed wing multi engine AIRBUS... A320...     2   182    NA Turbo...
# ... with abbreviated variable name 1manufacturer
```

Keys: controlling how the tables are matched

```
flights2 %>% left_join(planes, by = "tailnum")
```

```
# A tibble: 336,776 × 16
```

	year.x	month	day	hour	origin	dest	tailnum	carrier	year.y	type	manuf... ¹
	<int>	<int>	<int>	<dbl>	<chr>	<chr>	<chr>	<chr>	<int>	<chr>	<chr>
1	2013	1	1	5	EWB	IAH	N14228	UA	1999	Fixed w...	BOEING
2	2013	1	1	5	LGA	IAH	N24211	UA	1998	Fixed w...	BOEING
3	2013	1	1	5	JFK	MIA	N619AA	AA	1990	Fixed w...	BOEING
4	2013	1	1	5	JFK	BQN	N804JB	B6	2012	Fixed w...	AIRBUS
5	2013	1	1	6	LGA	ATL	N668DN	DL	1991	Fixed w...	BOEING
6	2013	1	1	5	EWB	ORD	N39463	UA	2012	Fixed w...	BOEING
7	2013	1	1	6	EWB	FLL	N516JB	B6	2000	Fixed w...	AIRBUS...
8	2013	1	1	6	LGA	IAD	N829AS	EV	1998	Fixed w...	CANADA...
9	2013	1	1	6	JFK	MCO	N593JB	B6	2004	Fixed w...	AIRBUS
10	2013	1	1	6	LGA	ORD	N3ALAA	AA	NA	<NA>	<NA>

```
# ... with 336,766 more rows, 5 more variables: model <chr>, engines <int>,
```

```
# seats <int>, speed <int>, engine <chr>, and abbreviated variable name
```

```
# 1manufacturer
```

Matching keys

```
flights2 %>% left_join(airports, c("origin" = "faa"))
# A tibble: 336,776 × 15
   year month   day hour origin dest tailnum carrier name    lat lon alt
  <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>
1  2013     1     1     5 EWR   IAH  N14228  UA    Newar... 40.7 -74.2  18
2  2013     1     1     5 LGA   IAH  N24211  UA    La Gu... 40.8 -73.9  22
3  2013     1     1     5 JFK   MIA  N619AA  AA    John ... 40.6 -73.8  13
4  2013     1     1     5 JFK   BQN  N804JB  B6    John ... 40.6 -73.8  13
5  2013     1     1     6 LGA   ATL  N668DN  DL    La Gu... 40.8 -73.9  22
6  2013     1     1     5 EWR   ORD  N39463  UA    Newar... 40.7 -74.2  18
7  2013     1     1     6 EWR   FLL  N516JB  B6    Newar... 40.7 -74.2  18
8  2013     1     1     6 LGA   IAD  N829AS  EV    La Gu... 40.8 -73.9  22
9  2013     1     1     6 JFK   MCO  N593JB  B6    John ... 40.6 -73.8  13
10 2013     1     1     6 LGA   ORD  N3ALAA  AA    La Gu... 40.8 -73.9  22
# ... with 336,766 more rows, and 3 more variables: tz <dbl>, dst <chr>,
#   tzone <chr>
```

inner_join()

```
df1 <- tibble(x = c(1, 2), y = 2:1)
df2 <- tibble(x = c(3, 1), a = 10, b = "a")
```

Table: df1

x	y
1	2
2	1

Table: df2

x	a	b
3	10	a
1	10	a

```
df1 %>% inner_join(df2)
```

x	y	a	b
1	2	10	a

left_join()

Table: df1

x	y
1	2
2	1

```
df1 %>% left_join(df2)
```

x	y	a	b
1	2	10	a
2	1	NA	NA

Table: df2

x	a	b
3	10	a
1	10	a

```
df2 %>% left_join(df1)
```

x	a	b	y
3	10	a	NA
1	10	a	2

right_join()

Table: df1

x	y
1	2
2	1

Table: df2

x	a	b
3	10	a
1	10	a

```
df1 %>% right_join(df2)
```

x	y	a	b
1	2	10	a
3	NA	10	a

```
df2 %>% right_join(df1)
```

x	a	b	y
1	10	a	2
2	NA	NA	1

Filtering joins

Filtering joins return a copy of the dataset that has been filtered, not augmented (as with mutating joins)

- `semi_join(x,y)` : keeps all observations in x that have a match in y.
- `anti_join(x,y)` : drops all observations in x that have a match in y.

most useful for diagnosing join mismatches

Another example

```
df1 <- tibble(x = c(1, 1, 3, 4), y = 1:4)
df2 <- tibble(x = c(1, 1, 2), z = c("a", "b", "a"))
```

Table: df1

x	y
1	1
1	2
3	3
4	4

Table: df2

x	z
1	a
1	b
2	a

semi_join()

Table: df1

x	y
1	1
1	2
3	3
4	4

```
df1 %>% semi_join(df2, by = "x")
```

x	y
1	1
1	2

Table: df2

x	z
1	a
1	b
2	a

```
df2 %>% semi_join(df1, by = "x")
```

x	z
1	a
1	b

anti_join()

Table: df1

x	y
1	1
1	2
3	3
4	4

```
df1 %>% anti_join(df2, by = "x")
```

x	y
3	3
4	4

Table: df2

x	z
1	a
1	b
2	a

```
df2 %>% anti_join(df1, by = "x")
```

x	z
2	a

Set Operations

These expect the x and y inputs to have the same variables, and treat the observations like sets:

- `intersect(x,y)`
 - will return only the rows that appear in both datasets
- `union(x,y)`
 - return every row that appears in one or more of the datasets
 - If a row appears multiple times union will only return it once
- `setdiff(x,y)`
 - will return the rows that appear in the first dataset but not the second

One more example

```
df1 <- tibble(x = 1:2, y = c(1L, 1L))  
df2 <- tibble(x = 1:2, y = 1:2)
```

Table: df1

x	y
1	1
2	1

Table: df2

x	y
1	1
2	2

Set operations

```
intersect(df1, df2))
```

x	y
1	1

```
union(df1, df2))
```

x	y
1	1
2	1
2	2

```
setdiff(df1, df2))
```

x	y
2	1

```
setdiff(df2, df1))
```

x	y
2	2

Group Activity 2

10:00



- Work on problems 3-5
- Ask me questions
- Any hw-related questions?