

More Regular Expressions

Stat 220

Bastola

February 01 2022

Quantifiers and Special Characters

Preceding characters are matched

- `*` = 0 or more
- `?` = 0 or 1
- `+` = 1 or more
- `{n}` = exactly n times

Matching character types

- `\\d` = digit
- `\\s` = white space
- `\\w` = word
- `\\t` = tab
- `\\n` = newline

Repetition using ?

```
aboutMe <- c("my SSN 536-76-9423 and my age is 55")
```

```
str_view_all(aboutMe, "\\s\\d?") # space followed by 0 or 1 digit
```

my SSN 536-76-9423 and my age is 55

Repitition using +

```
aboutMe <- c("my SSN 536-76-9423 and my age is 55")
```

```
str_view_all(aboutMe, "\\s\\d+") # space followed by 1 or more digits
```

my SSN 536-76-9423 and my age is 55

Repitition using *

```
aboutMe <- c("my SSN 536-76-9423 and my age is 55")
```

```
str_view_all(aboutMe, "\\s\\d*") # space followed by 0 or more digits
```

my SSN 536-76-9423 and my age is 55

More quantifiers

useful when you want to match a pattern a specific number of times

- $\{n, \}$ = n or more times
- $\{, m\}$ = at most m times
- $\{n, m\}$ = between n & m times

Alternatives

useful for matching patterns more flexibly

- `[abc]` = one of **a**, **b**, or **c**
- `[e-z]` = a letter from **e** to **z**
- `[^abc]` = anything other than **a**, **b**, or **c**

Recap: Extract Substrings

Using `str_sub()`

```
cc <- "Carleton College"  
str_sub(cc, start = 1, end = 8)  
[1] "Carleton"
```

```
str_sub(cc, 10) # end defaults to the last character  
[1] "College"
```

```
# match the elements of each vector for positions  
str_sub(cc, c(1, 10), c(8, 16))  
[1] "Carleton" "College"
```


Extract substrings

Can also extract the end..

```
cc <- "Carleton College"  
str_sub(cc, -4)  
[1] "lege"
```

```
str_sub(cc, -10, -4)  
[1] "on Coll"
```

Your Turn 1

Please git clone the repository on [advanced string manipulations](#) to your local folder.

Make changes to the code to ...

1. Isolate the last letter of every name
2. Create a logical variable that displays whether the last letter is one of "a", "e", "i", "o", "u", or "y".
3. Use a weighted mean to calculate the proportion of children whose name ends in a vowel by year
(see `?weighted.mean`)
4. and then display the results as a line plot.

04:00

str_count()

Tells you how many times a pattern appears in each entry of a string/character vector

```
days <- rep(c("Monday", "Tuesday", "Wednesday",  
              "Thursday", "Friday", "Saturday", "Sunday"), 2)
```

#number of times the letter "e" appears in each entry

```
days %>% str_count("e")  
[1] 0 1 2 0 0 0 0 0 1 2 0 0 0 0
```

number of times the letter "a" or "e" appears in each entry

```
days %>% str_count("[ae]")  
[1] 1 2 3 1 1 2 1 1 2 3 1 1 2 1
```

Your Turn 2

Determine how many baby names in 2017 contain "ders".

```
babynames %>%  
  filter(year == 2017) %>%  
  head(4)  
# A tibble: 4 × 5  
   year sex  name      n    prop  
  <dbl> <chr> <chr>   <int> <dbl>  
1  2017 F    Emma  19738 0.0105  
2  2017 F    Olivia 18632 0.00994  
3  2017 F    Ava    15902 0.00848  
4  2017 F    Isabella 15100 0.00805
```

02:00

str_extract()

pull all set of values matching the specified pattern

```
name_phone <- c("Moly Robins: 250-999-8878",  
               "Ali Duluth: 416-908-2044",  
               "Eli Mitchell: 204-192-9829",  
               "May Flowers: 250-209-7047")
```

```
str_extract(name_phone, "[:alpha:]*")  
[1] "Moly" "Ali"  "Eli"  "May"
```

```
str_extract(name_phone, "[:digit:]{3}-[:digit:]{3}-[:digit:]{4}")  
[1] "250-999-8878" "416-908-2044" "204-192-9829" "250-209-7047"
```

str_extract_all()

pull all set of values matching the specified pattern

```
name_phone <- c("Moly Robins: 250-999-8878",  
                "Ali Duluth: 416-908-2044",  
                "Eli Mitchell: 204-192-9829",  
                "May Flowers: 250-209-7047")
```

```
str_extract_all(name_phone, "[:alpha:]{2,}", simplify = TRUE)  
      [,1]      [,2]  
[1,] "Moly"    "Robins"  
[2,] "Ali"     "Duluth"  
[3,] "Eli"     "Mitchell"  
[4,] "May"     "Flowers"
```

Your Turn 3

Consider the 2017 babynames data

- Fill in the code to determine how many baby names in 2017 began with a vowel.
- Fill in the code to determine how many baby names in 2017 started and/or ended with a vowel.

04:00

Modify the case of a string

```
str_to_lower("BEAUTY is in the EYE of the BEHOLDER")  
[1] "beauty is in the eye of the beholder"
```

```
str_to_upper("one small step for man, one giant leap for mankind")  
[1] "ONE SMALL STEP FOR MAN, ONE GIANT LEAP FOR MANKIND"
```

```
str_to_title("Aspire to inspire before we expire")  
[1] "Aspire To Inspire Before We Expire"
```

```
str_to_sentence("everything you can imagine is real")  
[1] "Everything you can imagine is real"
```


Your Turn 4

Consider the 2017 babynames data

- how many baby names contain "illie`?
- how many baby names start with "M"?

04:00

Visualizing matches

```
str_view(name_phone, "[5-8]")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204-192-9829

May Flowers: 250-209-7047

```
str_view_all(name_phone, "[5-8]")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204-192-9829

May Flowers: 250-209-7047

Visualizing matches

```
str_view(name_phone, "[5|8]")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204-192-9829

May Flowers: 250-209-7047

```
str_view_all(name_phone, "[5-8|a-e]")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204-192-9829

May Flowers: 250-209-7047

Extract using repitition

```
str_view_all(name_phone, pattern = "[2-9][0-9]{2}-[0-9]{3}-[0-9]{4}")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204-192-9829

May Flowers: 250-209-7047

str_locate()

tell you the character positions the pattern characters are found in

```
days %>%  
  str_locate("day") %>% # start and end positions of the pattern "day"  
  head(n=7)  
      start end  
[1,]      4  6  
[2,]      5  7  
[3,]      7  9  
[4,]      6  8  
[5,]      4  6  
[6,]      6  8  
[7,]      4  6
```

str_flatten()

collapses a string vector into a single string (i.e. a character vector of length 1)

```
str_flatten(c("a_3", "d_54"), # vector to collapse
            collapse = " ") # insert in between
[1] "a_3 d_54"
```

str_glue()

allows one to interpolate strings and values that have been assigned to names in R

```
y <- Sys.Date() # current date
str_glue("today is {y}")
today is 2022-02-01
```

```
# base R equivalent
paste0("today is ", y)
[1] "today is 2022-02-01"
```

```
nm <- "Alex"
str_glue("Hi, my name is {nm}.")
Hi, my name is Alex.
```

```
a <- 5
str_glue("a = {a}")
a = 5
```

Lengths of strings

We can manage the lengths of strings using the following set of functions:

- `str_length()`
- `str_pad()`
- `str_trunc()`
- `str_trim()`

str_length()

tells you how many characters are in each entry of a character vector

```
gapminder %>% names()  
[1] "country" "continent" "year" "lifeExp" "pop" "gdpPercap"
```

length of each column names

```
gapminder %>% names() %>% str_length()  
[1] 7 9 4 7 3 9
```

str_pad()

standardizes the length of strings in a character vector by padding it on the left or right ends with a specified character (by default, a space)

```
gapminder %>% names() %>%  
  str_pad(width = 9, # minimum width of the string to fill/pad positions up to  
          side = "both", # side to insert the extra characters on  
          pad = "+")  
[1] "+country+" "continent" "++year+++" "+lifeExp+" "+++pop+++" "gdpPercap"
```

str_trunc()

standardizes string lengths by controlling the maximum width and truncating strings longer than this

```
gapminder %>% names() %>%  
  str_trunc(width = 7, #the maximum width to allow for strings  
            side= "right",  
            # entries which have been truncated will show this to indicate that  
            # something has been removed  
            ellipsis = "...")  
[1] "country" "cont..." "year"      "lifeExp" "pop"      "gdpP..."
```

str_trim()

removes empty spaces on the ends of a string

```
#add some whitespace
padded_names <- gapminder %>% names() %>% str_pad(12, "both")
padded_names
[1] "  country  " "continent" "  year  " "lifeExp" "  pop  "
[6] "gdpPercap"
```

```
#remove it
str_trim(padded_names)
[1] "country" "continent" "year" "lifeExp" "pop" "gdpPercap"
```

Disambiguate your regexps

```
name_phone <- c("Moly Robins: 250-999-8878",  
               "Ali Duluth: 416-908-2044",  
               "Eli Mitchell: 204.192.9829",  
               "May Flowers: 250.209.7047")
```

```
str_view_all(name_phone,  
             pattern = "([2-9][0-9]{2})[.-]([0-9]{3})[.-]([0-9]{4})")
```

Moly Robins: 250-999-8878

Ali Duluth: 416-908-2044

Eli Mitchell: 204.192.9829

May Flowers: 250.209.7047

Replacing strings

```
str_replace_all(name_phone,  
pattern = "([2-9][0-9]{2})[.-]([0-9]{3})[.-]([0-9]{4})",  
replacement = "XXX-XXX-XXXX"  
)  
[1] "Moly Robins: XXX-XXX-XXXX" "Ali Duluth: XXX-XXX-XXXX"  
[3] "Eli Mitchell: XXX-XXX-XXXX" "May Flowers: XXX-XXX-XXXX"
```

Duplicating Groups

Use escaped numbers (\1, \2, etc) to repeat a group based on position

Which numbers have the same 1st and 3rd digits?

```
phone_numbers <- c("515 111 2244", "310 549 6892", "474 234 7548")  
str_view(phone_numbers, "(\\d)\\d\\1")
```

515 111 2244

310 549 6892

474 234 7548

Your Turn 5

```
x <- c("scuttlebutt", "Stetson", "Scattter", "Scatter")  
str_detect(x, "^[Ss](.*)(t+)(.+) (t+)")  
[1]  TRUE  TRUE  TRUE FALSE
```

The regular expression `"^[Ss](.*)(t+)(.+) (t+)"` detects "scuttlebutt", "Stetson", and "Scattter", but not "Scatter." Why?

05:00