


ANOVA for MLR

Stat 230

April 22 2022

Overview




Today:

- ANOVA for MLR
- MLR Inference
- R^2 adjusted
- ANOVA F-test

Penguins example

$$\begin{aligned}\mu_{\text{mass} | x} = & \beta_0 + \beta_1 \text{bill} + \beta_2 \text{speciesChinstrap} + \beta_3 \text{speciesGentoo} \\ & + \beta_4 \text{bill} \times \text{speciesChinstrap} + \beta_5 \text{bill} \times \text{speciesGentoo}\end{aligned}$$



- We will test whether the effect of bill length on mass depends on species:

$$H_0 : \beta_4 = \beta_5 = 0$$

Comparing models

Is the effect of bill length the same for Gentoo and Adelie (baseline)?

$$H_0 : \beta_5 = 0$$

$$H_A : \beta_5 \neq 0$$

Equivalent test:

$$H_0 : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo} \\ + \beta_4 \text{ bill} \times \text{speciesChinstrap}$$

$$H_A : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo} \\ + \beta_4 \text{ bill} \times \text{speciesChinstrap} + \beta_5 \text{ bill} \times \text{speciesGentoo}$$

Individual t-tests: Is this term significant assuming all others are in the model?

Comparing models

Does the effect of bill length on mass depends on species?

$$H_0 : \beta_4 = \beta_5 = 0$$

$$H_A : \text{at least one } \beta_4, \beta_5 \neq 0$$

Equivalent test:

$$H_0 : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo}$$

$$H_A : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo} \\ + \beta_4 \text{ bill} \times \text{speciesChinstrap} + \beta_5 \text{ bill} \times \text{speciesGentoo}$$

Our new test ("F-test"): Is at least one of these terms significant assuming all others are in the model?

Comparing models

$$H_0 : \beta_4 = \beta_5 = 0$$

$$H_A : \text{at least one } \beta_4, \beta_5 \neq 0$$

We can no longer use

$$t = \frac{\hat{\beta} - 0}{SE(\hat{\beta})}$$

We will use Analysis of Variance (ANOVA)

- **idea:** how much more explanatory power does the bigger model H_A have compared to the smaller model H_0 ?

ANOVA for MLR (basically same as SLR!)

$$SST = SSreg + SSR$$

- SST: Total variation Total sum of squares

$$SST = \sum_{i=1}^n (y_i - \bar{y})^2 = (n - 1)s_y^2$$

- SSR: Unexplained variation Residual sum of squares

$$SSR = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = (n - p - 1)\hat{\sigma}^2$$

- SSreg: Explained variation Regression sum of squares

$$SSreg = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$$


R-squared

R-squared (coefficient of determination) measures the proportion of variability observed in the response Y which can be explained by the regression of \mathbf{Y} on x_1, \dots, x_p .

$$R^2 = \frac{SS_{reg}}{SST} = 1 - \frac{SSR}{SST}$$

- About 35.4% of the variation in mass is explained by the regression of mass on bill length.
- About 78.3% of the variation in mass is explained by the regression of mass on bill length and species.
- About 78.8% of the variation in mass is explained by the regression of mass on bill length, species and their interaction.

Penguins: ANOVA for 3 models



$$R^2 = \frac{SS_{reg}}{SST} = 1 - \frac{SSR}{SST}$$

- SS_{reg} (explained) is largest for the biggest model
- SSR (unexplained) is smallest for the biggest model

ANOVA for MLR

$$SST = SSreg + SSR$$

- **Mathematically:** bigger models have bigger SSreg:

$$SSreg(x_1, x_2, x_1 : x_2) > SSreg(x_1, x_2) > SSreg(x_1)$$

- **Mathematically:** bigger models have smaller SSR:

$$SSR(x_1, x_2, x_1 : x_2) < SSR(x_1, x_2) < SSR(x_1)$$

- Bigger models have more explanatory power (higher SSreg) and larger R^2 than a smaller model
- Assume the smaller model is a nested version of the larger
- But is the bigger model "worth it" if it is more complex and only has slightly more SSreg explained?

ANOVA for MLR

Mean square values divide SS by their degrees of freedom

- Degrees of freedom (basically) measure how much stuff we've added up in the SS
- **MST:** Average total variation Mean Square of Total

$$MST = \frac{SST}{n - 1} = s_y^2 = \text{sample variance of Y}$$

- **MSR:** Average unexplained variation residual mean square

$$MSR = \frac{SSR}{n - (p + 1)} = \hat{\sigma}^2 = \text{estimated model variance}$$

- **MSreg:** Explained variation Regression mean square

$$MSreg = \frac{SSreg}{p}$$

ANOVA for MLR

$$SST = SSreg + SSR$$

- **Mathematically:** mean square residual and regression values won't equal the total

$$MST \neq MSreg + MSR$$

- **Mathematically:** mean square regression values won't be ordered based on their size

$$MSreg(x_1, x_2, x_1 : x_2)? MSreg(x_1, x_2)? MSreg(x_1)$$

- **Mathematically:** mean square residual values won't be ordered based on their size

$$MSR(x_1, x_2, x_1 : x_2)? MSR(x_1, x_2)? MSR(x_1)$$

Adjusted R-squared

Adjusted R^2 is similar to R^2 , but it uses mean square values to account for the size of a model.

$$R_a^2 = R_{\text{adjust}}^2 = 1 - \frac{MSR}{MST}$$

- R_a^2 could actually increase if unneeded predictors are removed from a model.
- R^2 would never increase when predictors are removed.

Penguins: Adjusted R^2 for 3 models

```
penguins <- penguins %>% tidyr::drop_na(bill_length_mm, body_mass_g, species) # drop missing values
peng_bill_lm <- lm(body_mass_g ~ bill_length_mm, data = penguins) # simple linear regression model
peng_nointeraction_lm <- lm(body_mass_g ~ bill_length_mm + species, data = penguins) # parallel lines model
peng_interaction_lm <- lm(body_mass_g ~ bill_length_mm*species, data = penguins) # interaction model
```

```
get_regression_summaries(peng_bill_lm) %>%
  select(r_squared, adj_r_squared) # just want these two
# A tibble: 1 × 2
  r_squared adj_r_squared
    <dbl>      <dbl>
1    0.354      0.352
```

```
get_regression_summaries(peng_nointeraction_lm) %>%
  select(r_squared, adj_r_squared)
# A tibble: 1 × 2
  r_squared adj_r_squared
    <dbl>      <dbl>
1    0.783      0.781
```

```
get_regression_summaries(peng_nointeraction_lm) %>%
  select(r_squared, adj_r_squared)
# A tibble: 1 × 2
  r_squared adj_r_squared
    <dbl>      <dbl>
1    0.783      0.781
```

```
glimpse(penguins)
```

```
Rows: 342
Columns: 8
$ species      <fct> Adelie, Adelie,
$ island       <fct> Torgersen, Torge
$ bill_length_mm <dbl> 39.1, 39.5, 40.3
$ bill_depth_mm <dbl> 18.7, 17.4, 18.6
$ flipper_length_mm <int> 181, 186, 195, 1
$ body_mass_g   <int> 3750, 3800, 3250
$ sex          <fct> male, female, fe
$ year         <int> 2007, 2007, 2007
```

Penguins: Adjusted R^2 for 3 models

- Removing the interaction terms results in a small drop in adjusted R^2 (0.785 to 0.781).
- The explanatory power of both the larger interaction model and smaller no-interaction models looks similar.
- **Next:** let's verify this observation with an ANOVA F-test

Comparing models: ANOVA F-test

- **Hypotheses:**

$$H_0 : \mu_{Y|x} = \text{smaller, nested model}$$

$$H_A : \mu_{Y|x} = \text{larger model}$$

where the nested model is arrived at by setting one or more β_j in the larger model equal to 0

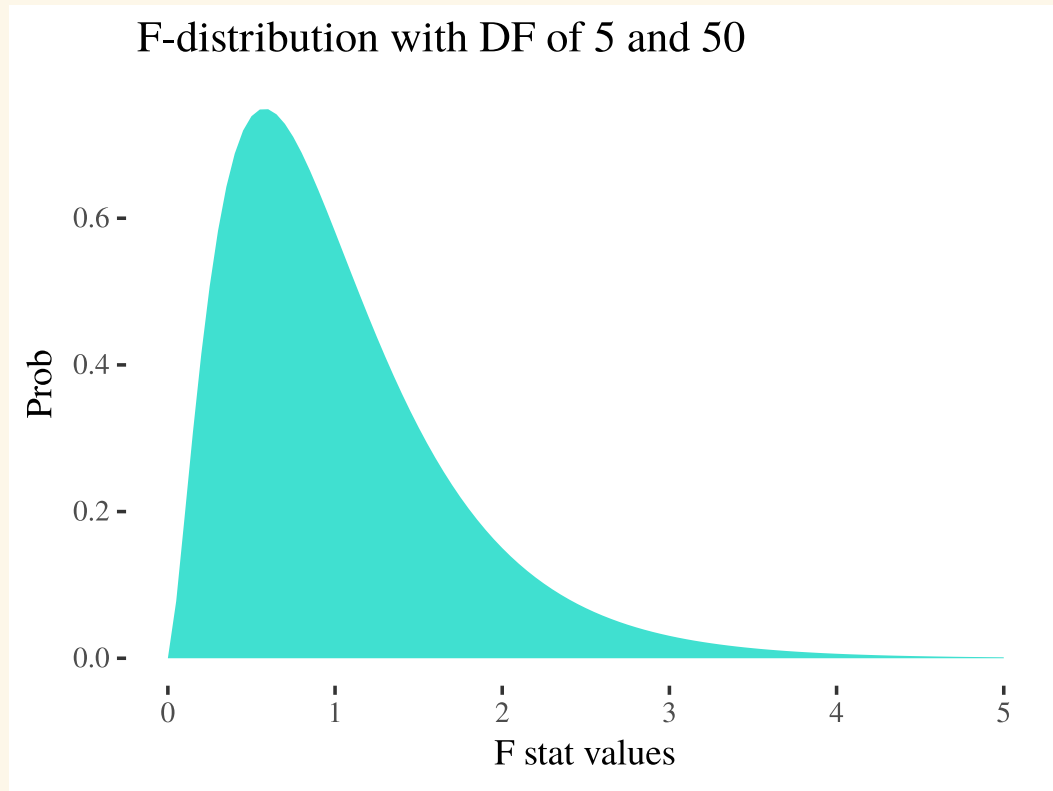
- **Test stat:**

$$F = \frac{\text{extraSS} / (\# \text{ terms tested})}{MSR_{\text{larger}}}$$

where extra SS is the difference in explained SS between the larger H_A and smaller H_0 models

$$\begin{aligned} \text{extraSS} &= SS_{\text{reg}}(\text{larger model}) - SS_{\text{reg}}(\text{smaller model}) \\ &= SSR(\text{smaller model}) - SSR(\text{larger model}) \end{aligned}$$

Comparing models: ANOVA F-test



- **P-value:** Probability of getting an T test stat larger than one observed

$$\text{p-value} = P(F > f_{\text{observed}})$$

using the F-distribution with degrees of freedom equal to the # of terms tested (numerator) and $n - (p + 1)$ from the larger H_A model (denominator)

ANOVA F-test in R

R-code

```
anova(smaller_lm, larger_lm)
```

Let's compare the interaction and parallel lines models in penguins dataset:

$$H_0 : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo}$$

$$H_A : \mu_{\text{mass} | x} = \beta_0 + \beta_1 \text{ bill} + \beta_2 \text{ speciesChinstrap} + \beta_3 \text{ speciesGentoo} \\ + \beta_4 \text{ bill} \times \text{speciesChinstrap} + \beta_5 \text{ bill} \times \text{speciesGentoo}$$

- Same as testing β_4 and β_5 in the larger model

$$H_0 : \beta_4 = \beta_5 = 0$$

$$H_A : \text{at least one } \beta_4, \beta_5 \neq 0$$

ANOVA F-test in R

```
anova(peng_nointeraction_lm , peng_interaction_lm)
```

Analysis of Variance Table

Model 1: `body_mass_g ~ bill_length_mm + species`

Model 2: `body_mass_g ~ bill_length_mm * species`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	338	47613707				
2	336	46447006	2	1166702	4.22	0.01549 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

ANOVA F-test in R

```
anova(peng_nointeraction_lm , peng_interaction_lm)
```

Analysis of Variance Table

Model 1: `body_mass_g ~ bill_length_mm + species`

Model 2: `body_mass_g ~ bill_length_mm * species`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	338	47613707				
2	336	46447006	2	1166702	4.22	0.01549 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Res.Df column

- model DF for smaller = $342 - 4 = 338$
- model DF for larger = $342 - 6 = 336$

RSS Df column

- RSS(smaller)= 47,613,707
- RSS(larger)= 46,447,006

ANOVA F-test in R

```
anova(peng_nointeraction_lm , peng_interaction_lm)
```

Analysis of Variance Table

```
Model 1: body_mass_g ~ bill_length_mm + species
```

```
Model 2: body_mass_g ~ bill_length_mm * species
```

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	338	47613707				
2	336	46447006	2	1166702	4.22	0.01549 *

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Df column

- number of terms tested
 $= 6 - 4 = 338 - 336 = 2$

Sum of Sq column

- Extra SS for this test = $\text{RSS}(\text{smaller}) - \text{RSS}(\text{larger}) = 1,166,702$

ANOVA F-test in R

```
anova(peng_nointeraction_lm , peng_interaction_lm)
```

Analysis of Variance Table

Model 1: body_mass_g ~ bill_length_mm + species

Model 2: body_mass_g ~ bill_length_mm * species

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	338	47613707				
2	336	46447006	2	1166702	4.22	0.01549 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


F column

- **F** test stat = $F = \frac{1,166,702/2}{46,447,006/336} = 4.22$

Pr(> F) column

- p-value for this test
= $P(F_{2,336} > 4.22) = 1 - \text{pf}(4.22, 2, 336)$
= 0.01549

Penguins wrapup



The slopes are different enough to warrant an interaction model

Penguins wrapup

```
summary(peng_interaction_lm)$coefficients # simple way of getting summaries
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.88299	443.17604	0.07871136	9.373091e-01
bill_length_mm	94.49982	11.39794	8.29095491	2.733899e-15
speciesChinstrap	811.26034	799.80552	1.01432201	3.111590e-01
speciesGentoo	-158.71092	683.19141	-0.23230813	8.164401e-01
bill_length_mm:speciesChinstrap	-35.38208	17.74666	-1.99373192	4.698940e-02
bill_length_mm:speciesGentoo	14.95935	15.78642	0.94760841	3.440099e-01

Penguins wrapup

```
summary(peng_interaction_lm)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.88299	443.17604	0.07871136	9.373091e-01
bill_length_mm	94.49982	11.39794	8.29095491	2.733899e-15
speciesChinstrap	811.26034	799.80552	1.01432201	3.111590e-01
speciesGentoo	-158.71092	683.19141	-0.23230813	8.164401e-01
bill_length_mm:speciesChinstrap	-35.38208	17.74666	-1.99373192	4.698940e-02
bill_length_mm:speciesGentoo	14.95935	15.78642	0.94760841	3.440099e-01

- The difference in slopes (effect of bill length on mass) between Adelie (baseline) and Gentoo isn't statistically significance ($t = 0.948$, $df = 336$, $p = 0.344$)

Penguins wrapup

```
summary(peng_interaction_lm)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.88299	443.17604	0.07871136	9.373091e-01
bill_length_mm	94.49982	11.39794	8.29095491	2.733899e-15
speciesChinstrap	811.26034	799.80552	1.01432201	3.111590e-01
speciesGentoo	-158.71092	683.19141	-0.23230813	8.164401e-01
bill_length_mm:speciesChinstrap	-35.38208	17.74666	-1.99373192	4.698940e-02
bill_length_mm:speciesGentoo	14.95935	15.78642	0.94760841	3.440099e-01

- The difference in slopes (effect of bill length on mass) between Adelie (baseline) and Chinstrap is weakly statistically significance ($t = -1.994$, $df = 336$, $p = 0.047$).

Penguins wrapup

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.88299	443.17604	0.07871136	9.373091e-01
bill_length_mm	94.49982	11.39794	8.29095491	2.733899e-15
speciesChinstrap	811.26034	799.80552	1.01432201	3.111590e-01
speciesGentoo	-158.71092	683.19141	-0.23230813	8.164401e-01
bill_length_mm:speciesChinstrap	-35.38208	17.74666	-1.99373192	4.698940e-02
bill_length_mm:speciesGentoo	14.95935	15.78642	0.94760841	3.440099e-01

$$\begin{aligned}\mu_{\text{mass} | x} = & \beta_0 + \beta_1 \text{bill} + \beta_2 \text{speciesChinstrap} + \beta_3 \text{speciesGentoo} \\ & + \beta_4 \text{bill} \times \text{speciesChinstrap} + \beta_5 \text{bill} \times \text{speciesGentoo}\end{aligned}$$

- To compare Chinstrap and Gentoo, either use the test

$$H_0 : \beta_1 + \beta_4 = \beta_1 + \beta_5 \Rightarrow \beta_4 - \beta_5 = 0$$

Penguins wrapup

- OR! Relevel the species variable to make Chinstrap the baseline - it's slope seem most different (flatter) compared to the other two species

$$\begin{aligned}\mu_{\text{mass} | x} = & \beta_0^* + \beta_1^* \text{bill} + \beta_2^* \text{speciesAdelie} + \beta_3^* \text{speciesGentoo} \\ & + \beta_4^* \text{bill} \times \text{speciesAdelie} + \beta_5^* \text{bill} \times \text{speciesGentoo}\end{aligned}$$

- Same variables but changing up the indicator variables changes our parameter interpretation
- that's why we used β^* here (not the same values as β in the previous model)
- but won't change an inference conclusions or effect estimation!!

Penguins wrapup

- OR! Relevel the species variable to make Chinstrap the baseline
- Use `fct_relevel()` from the `forcats` package:

```
library(forcats)
# just need to specify the first level
penguins$new_species <- fct_relevel(penguins$species, "Chinstrap")
```

Check order using a quick table command:

```
table(penguins$species)
```

Adelie	Chinstrap	Gentoo
151	68	123

```
table(penguins$new_species)
```

Chinstrap	Adelie	Gentoo
68	151	123

Penguins wrapup

```
peng_interaction_lm_CS <- lm(body_mass_g ~ bill_length_mm * new_species, data = penguins)
summary(peng_interaction_lm_CS)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	846.14333	665.79566	1.270875	2.046521e-01
bill_length_mm	59.11773	13.60261	4.346059	1.838939e-05
new_speciesAdelie	-811.26034	799.80552	-1.014322	3.111590e-01
new_speciesGentoo	-969.97126	844.76586	-1.148213	2.516974e-01
bill_length_mm:new_speciesAdelie	35.38208	17.74666	1.993732	4.698940e-02
bill_length_mm:new_speciesGentoo	50.34143	17.44503	2.885718	4.157658e-03

Penguins wrapup

```
peng_interaction_lm_CS <- lm(body_mass_g ~ bill_length_mm * new_species, data = penguins)
summary(peng_interaction_lm_CS)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	846.14333	665.79566	1.270875	2.046521e-01
bill_length_mm	59.11773	13.60261	4.346059	1.838939e-05
new_speciesAdelie	-811.26034	799.80552	-1.014322	3.111590e-01
new_speciesGentoo	-969.97126	844.76586	-1.148213	2.516974e-01
bill_length_mm:new_speciesAdelie	35.38208	17.74666	1.993732	4.698940e-02
bill_length_mm:new_speciesGentoo	50.34143	17.44503	2.885718	4.157658e-03

- The difference in slopes (effect of bill length on mass) between Chinstrap (baseline) and Adelie is weakly statistically significance ($t = 1.994$, $df = 336$, $p = 0.047$).

Penguins wrapup

```
summary(peng_interaction_lm_CS)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	846.14333	665.79566	1.270875	2.046521e-01
bill_length_mm	59.11773	13.60261	4.346059	1.838939e-05
new_speciesAdelie	-811.26034	799.80552	-1.014322	3.111590e-01
new_speciesGentoo	-969.97126	844.76586	-1.148213	2.516974e-01
bill_length_mm:new_speciesAdelie	35.38208	17.74666	1.993732	4.698940e-02
bill_length_mm:new_speciesGentoo	50.34143	17.44503	2.885718	4.157658e-03

- The difference in slopes (effect of bill length on mass) between Chinstrap (baseline) and Gentoo is statistically significant ($t = 2.886$, $df = 336$, $p = 0.00416$).